

# Unit 12: Single Cycle Processors

CSE 220: System Fundamentals I

Stony Brook University

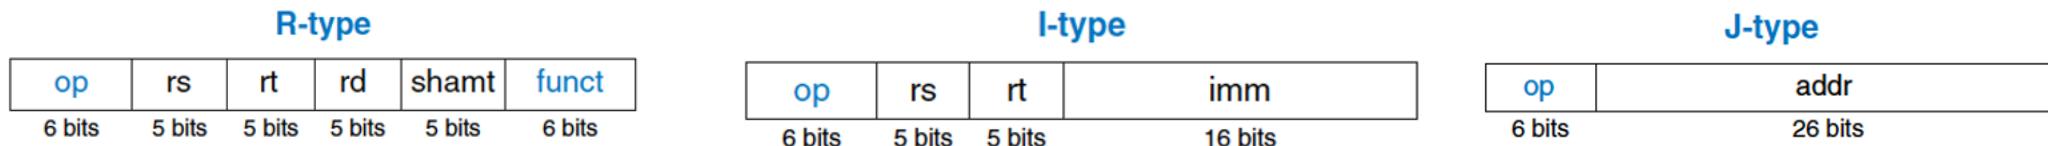
Joydeep Mitra

# Microarchitecture

- Microarchitecture is the connection between logic and architecture – a specific arrangement of registers, ALUs, memories, and other logic blocks.
- A particular architecture (e.g., MIPS) may have numerous microarchitectures, each with different trade-offs of performance, cost, and complexity.
- We will study two microarchitectures for the MIPS processor architecture
  - Single Cycle – executes an entire instruction in one cycle
  - Pipelining – executes several instructions simultaneously with the goal of improving throughput

# Architectural State And Instruction Set

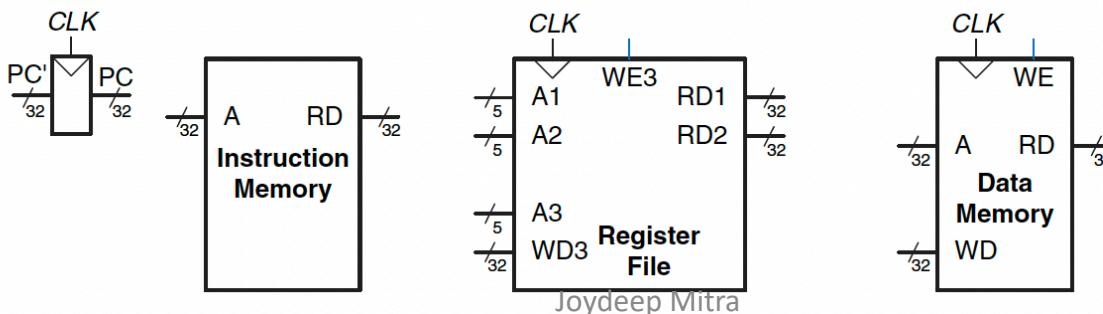
- Any computer architecture is defined by its architectural state and instruction set
- Recall the architectural state of MIPS is the program counter (keeps track next instruction to execute) and the 32 registers
- Also, recall the MIPS instruction formats – R-Type, I-Type, and J-Type



- A microarchitecture must remember all of this state
- Based on the current architectural state, a processor executes a particular instruction to produce a new state

# Design Process

- We will divide a microarchitecture into **datapath** and **control**.
  - The datapath operates on words of data and contains structures such as memories, registers, ALUs, and multiplexers
  - The control unit receives the current instruction and tells the datapath how to execute it.
- It is convenient to first define the hardware that will contain the architectural state.
  - **Program State** is 32-bit register. PC points to current instruction and PC' indicates next instruction address.
  - **Instruction Memory** has a single read port. It takes a 32-bit instruction address input, A, and reads the 32-bit instruction from that address onto the read data output, RD.
  - **Register File** has ports, A1, A2, and A3, which take 5-bit addresses. Data from A1 and A2 are read onto RD1 and RD2. Data from WD3 is written to A3. WE3 is a write enable signal, which if enabled (set to 1) allows writing to register on the rising edge of the clock.
  - **Data Memory** has a single read/write port. If WE is enabled, i.e., set to 1, then it writes data WD into A; otherwise, reads from address A onto RD

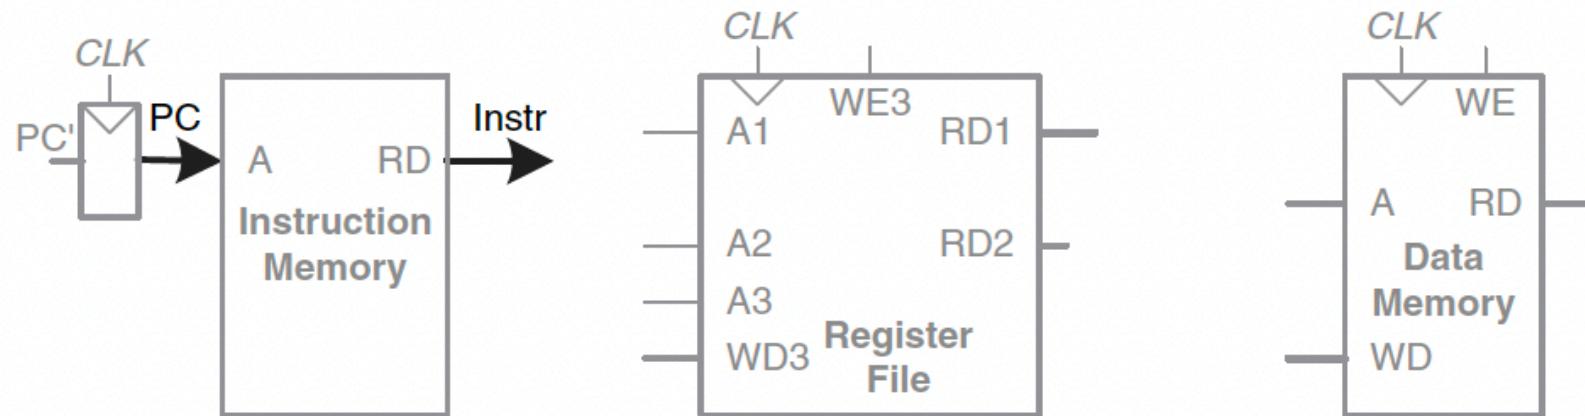


# Single Cycle Data Processor

- We will learn how to design a single cycle processor by first looking at the `lw` instruction
- Then we will generalize the design to include other instructions
- To keep things simple, we will focus on:
  - R-type arithmetic/logic instructions: `add`, `sub`, `and`, `or`, `slt`
  - Memory instructions: `lw`, `sw`
  - Branches: `beq`

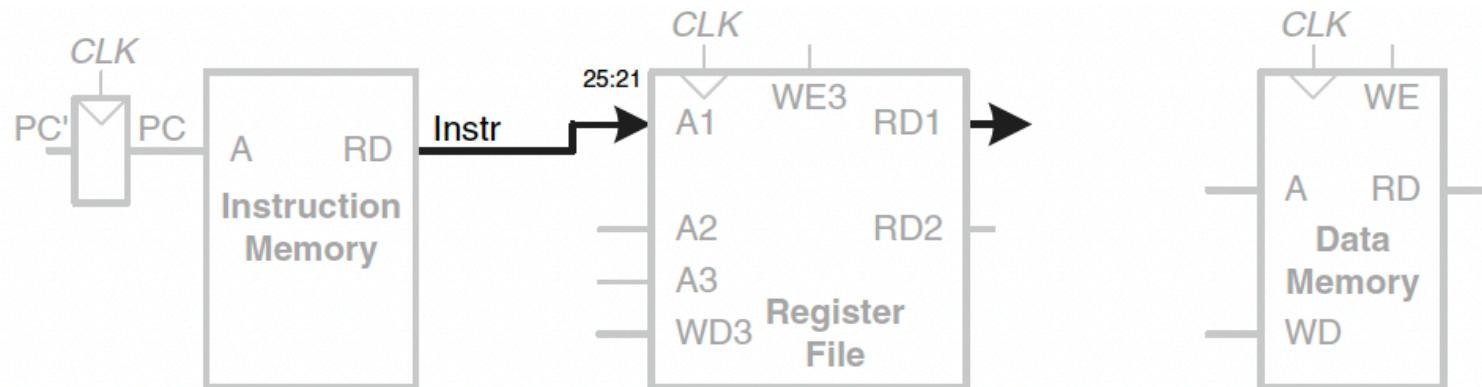
# Single Cycle Data Processor: $lw$

- The first step is to fetch an instruction from instruction memory



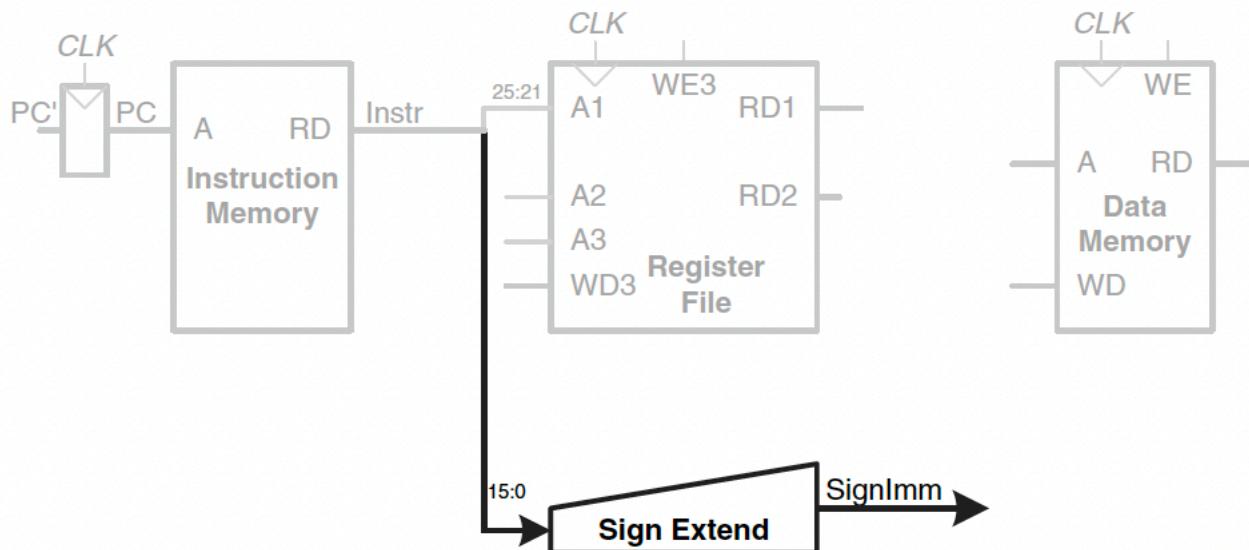
# Single Cycle Data Processor: $lw$

- The next step is to read the source register containing the base address
- The register is specified in the  $rs$  field of the instruction,  $\text{Instr}_{25:21}$
- Bits  $\text{Instr}_{25:21}$  are connected to A1 of register file
- The register file reads the register value onto RD1



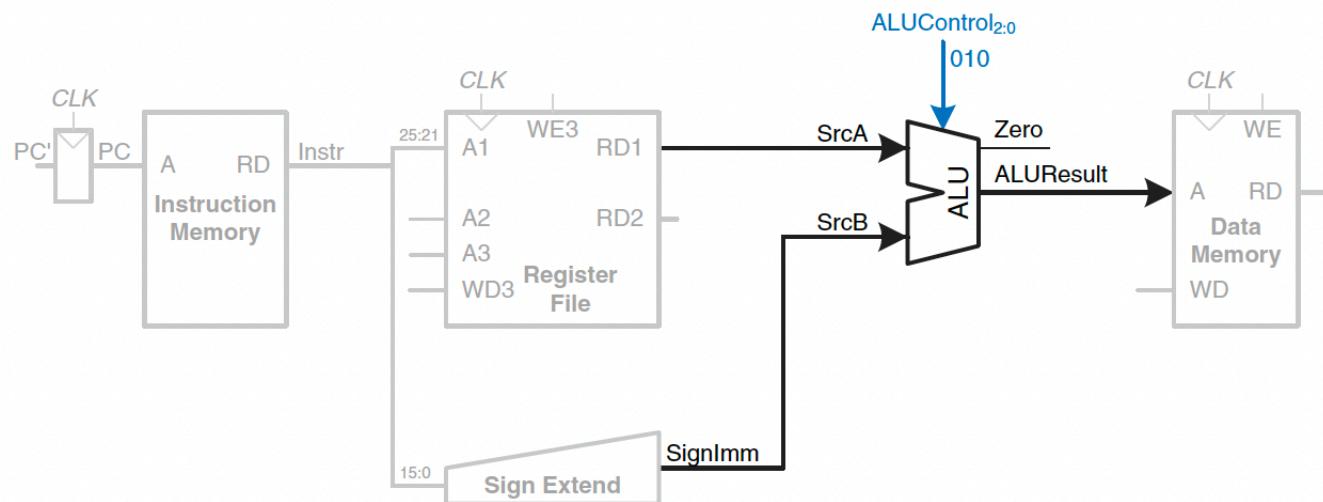
# Single Cycle Data Processor: `lw`

- The `lw` instruction also needs an offset, which is added to the base address
- The offset is stored in the immediate field of the instruction,  $\text{Instr}_{15:0}$
- The offset must be sign-extended since it can be positive or negative



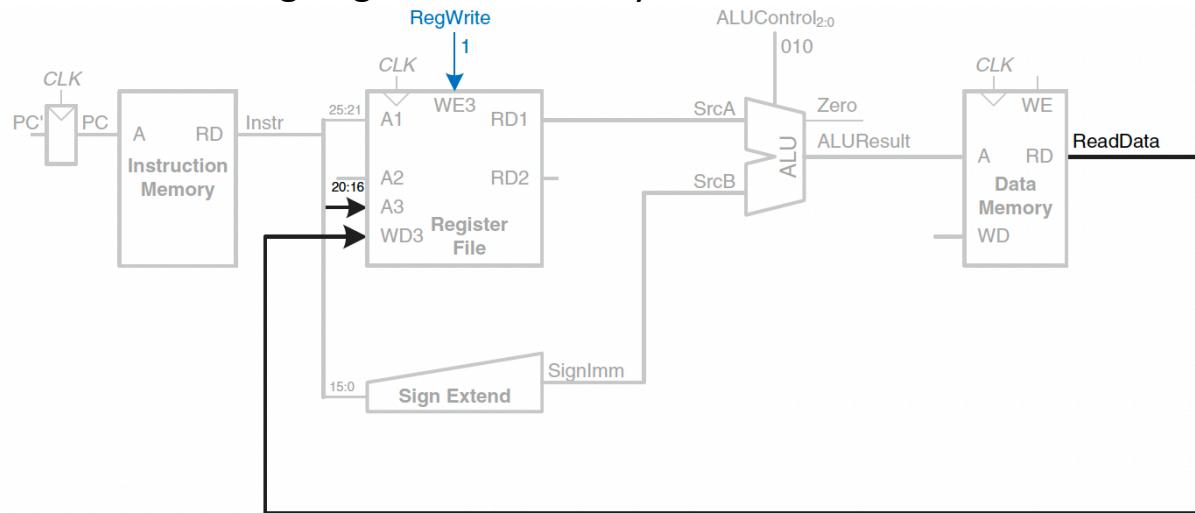
# Single Cycle Data Processor: $l_w$

- Next, we introduce an ALU to add the base address ( $SrcA$ ) and the offset ( $SrcB$ )
  - $SrcA$  comes from register file, and  $SrcB$  comes from sign-extended immediate
- The ALU can perform various operations based on the 3-bit  $ALUControl$  (assume it is 010 for *add* instruction for now)
- The ALU produces a 32-bit  $ALUResult$  and a *Zero* flag, that indicates if  $ALUResult == 0$
- $ALUResult$  is sent to the data memory as the address of the load instruction



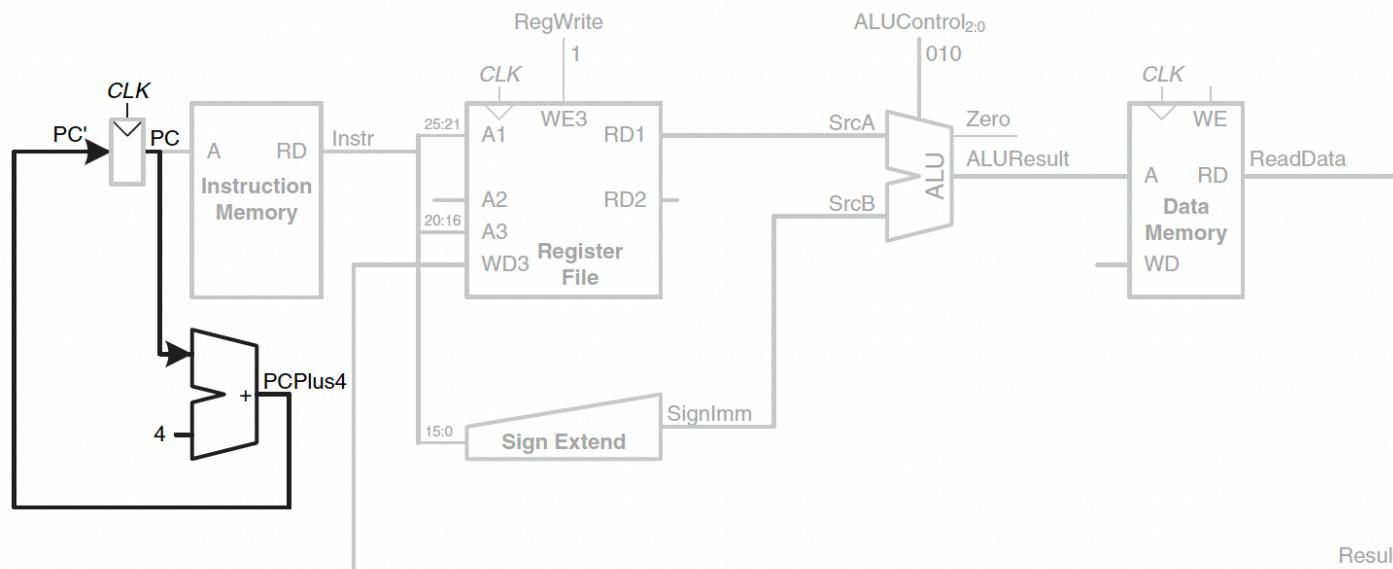
# Single Cycle Data Processor: **lw**

- Data from port A in *data memory* is read onto bus RD and connected back to destination register in the *register file* at the end of the cycle
- The destination register is the `rt` field in the **lw** instruction. Hence, port A3 is connected to bits  $\text{Instr}_{20:16}$
- The signal `WE3` (connected to `RegWrite`) must be set to 1 so that data in `WD3` can be written to destination register
  - Write takes place on the rising edge of the clock cycle



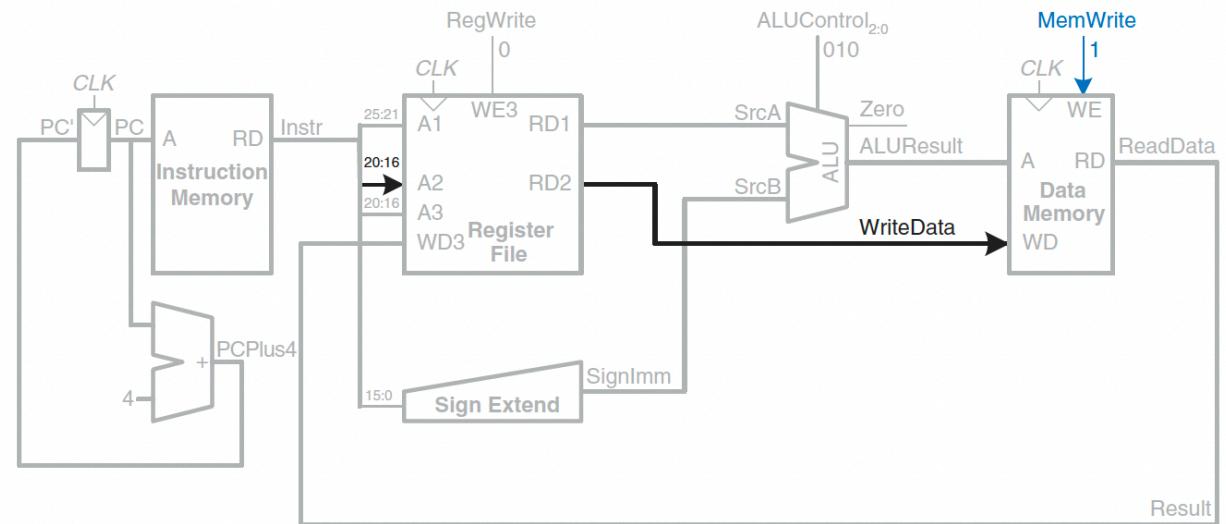
# Single Cycle Data Processor: $l_w$

- While  $l_w$  is executing, the processor must determine the next instruction. How?
- Increment the program counter by 4 (since instructions are 32-bits = 4 bytes)
- We introduce an additional adder
- $PC + 4$  is written to program counter on the rising edge of the clock cycle



# Single Cycle Data Processor: sw

- We will extend the datapath with `sw` instruction.
- The `sw` instruction reads from a register and stores in address + offset.
- Address calculation ALU already exists.
- From instruction format, we know the register field is `rt` (bits  $\text{Instr}_{20:16}$ ).
- These bits are connected to A2 in the register file and are read out of the RD2 port.
- *WriteData* from RD2 is then connected to WD port of Data Memory, which is then written to address A on the rising edge of the clock when *MemWrite* is set to 1.

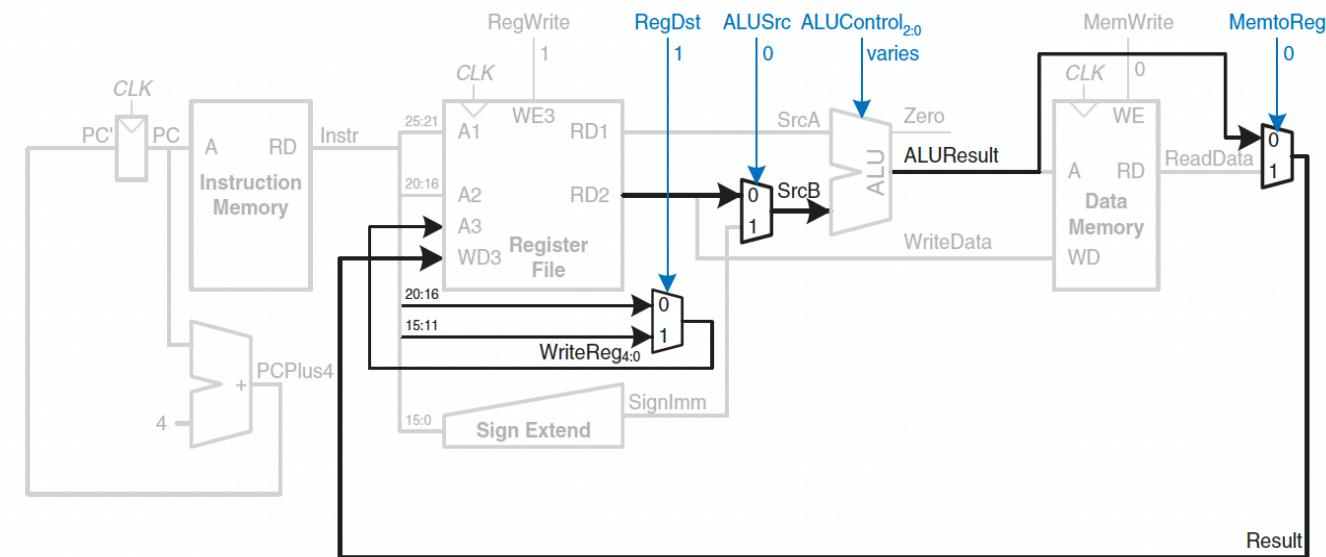


# Single Cycle Data Processor: RType

- We will extend the datapath with R-Type instructions – add, sub, and, or, and slt
- These instructions have a common pattern
  - Read two registers from register file
  - Perform some ALU operation on them
  - Write result back to a 3<sup>rd</sup> register file
- They can be handled with the same hardware, except for the ALU operation that can be handled with different *ALUControl* signals

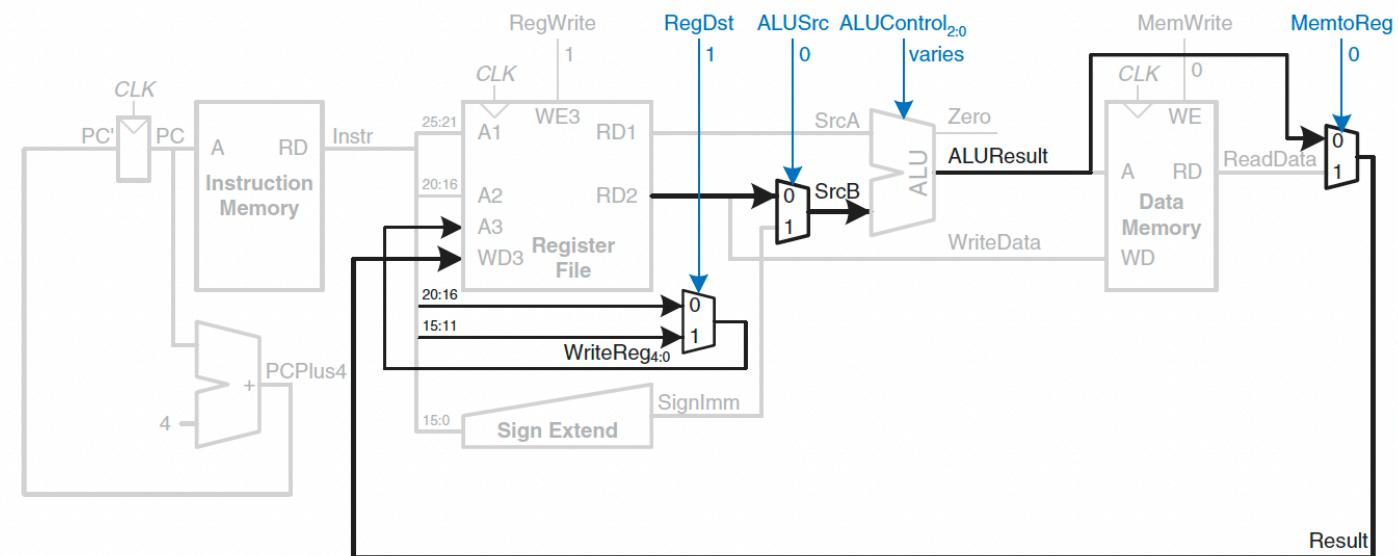
# Single Cycle Data Processor: RType

- The ALU operation needs two operands,  $SrcA$  and  $SrcB$ , from the register file
- But previously the 2<sup>nd</sup> operand to the ALU was always the sign-extended immediate
- Clearly, we need a multiplexer to choose between the  $RD2$  and  $SignImm$
- The control/select signal for the multiplexer is  $ALUSrc$ , which is 0 for R-Type instructions and 1 for  $lw$  and  $sw$



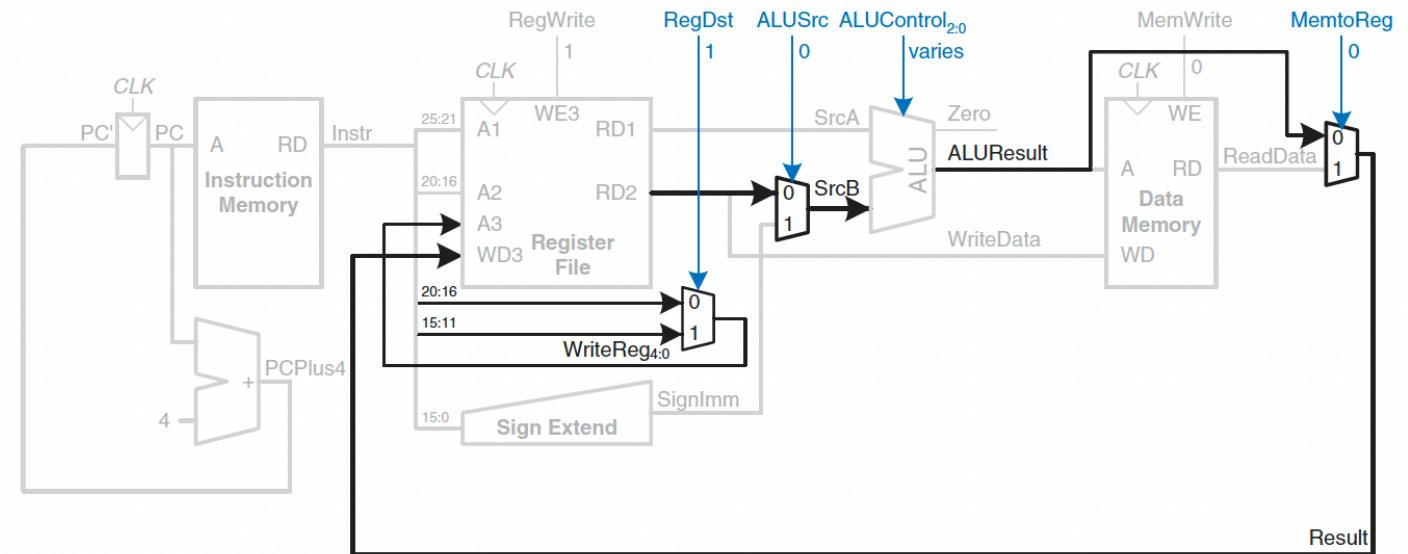
# Single Cycle Data Processor: RType

- *ALUResult* is written to the register file
- But *ReadData* from Data Memory is also written to register file
- We need another multiplexer to select whether *ALUResult* or *ReadData* will be written to register file
- This is controlled by *MemtoReg*, which is 0 for R-Type and 1 for the *lw* instruction
- What about *MemtoReg* for *sw* instruction?
  - Don't care as *sw* does not write to register file



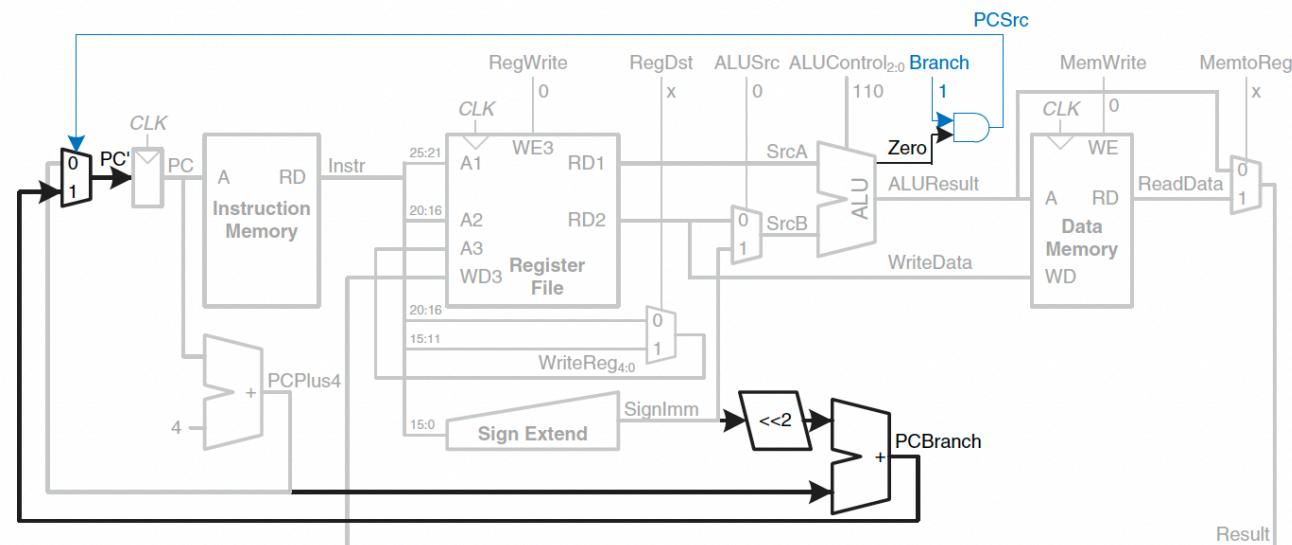
# Single Cycle Data Processor: RType

- In R-type instructions the `rd` field (bits  $\text{Instr}_{15:11}$ ) has the register address where the result will be written
- But in register file, the register address is always bits  $\text{Instr}_{20:16}$
- Hence, we need another multiplexer with  $\text{RegDst}$  as control to select between `rd` for R-type and `rt` for `lw` instruction
- $\text{RegDst}$  is 0 for R-type and 1 for `lw`



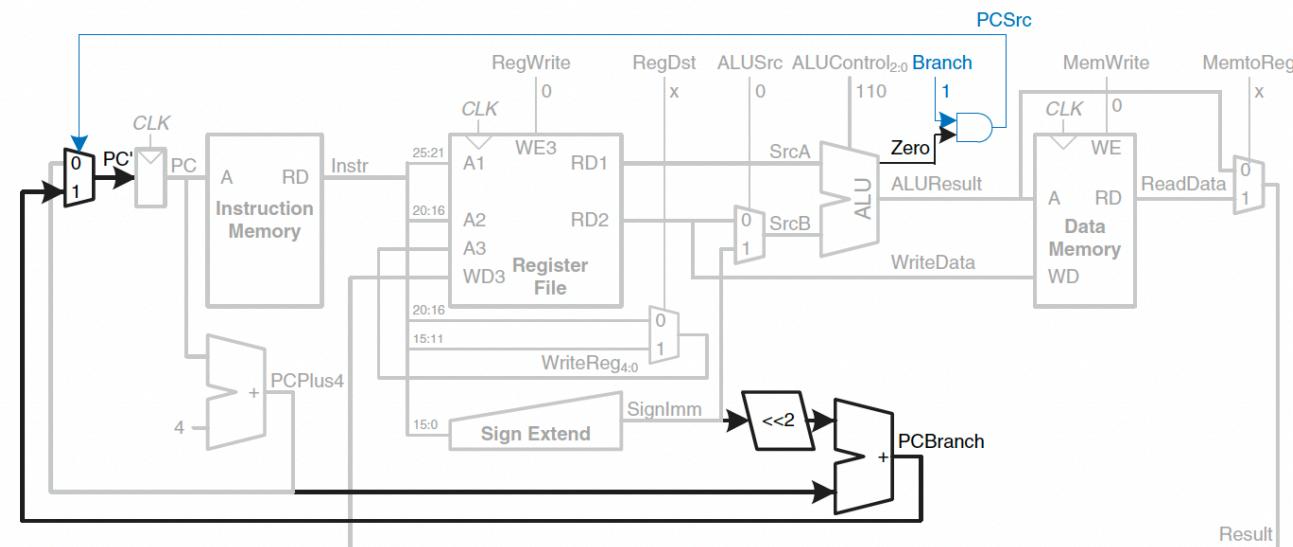
# Single Cycle Data Processor: beq

- We extend the datapath with the **beq** instruction
- **beq** compares two registers; if equal branch is taken else next instruction is executed
- The actual comparison involves  $\text{SrcA} - \text{SrcB}$ 
  - if *ALUResult* is 0 (Zero flag set to 1), then equal



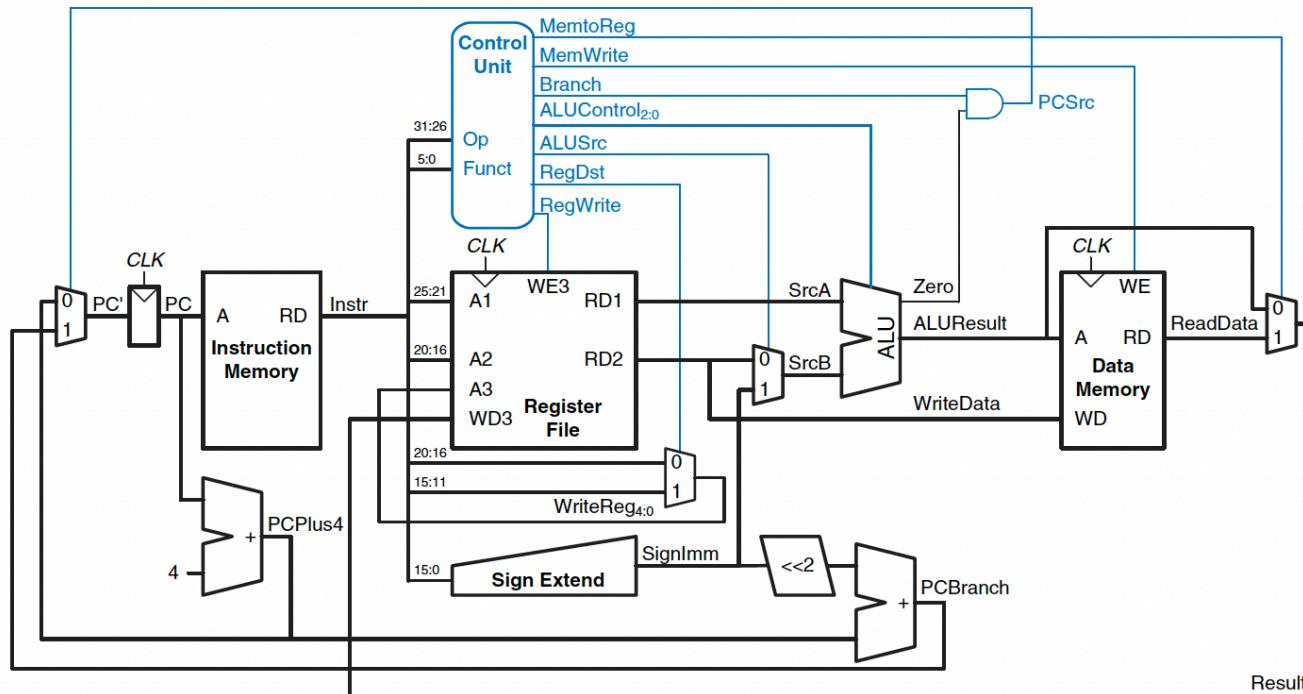
# Single Cycle Data Processor: beq

- A branch is taken by adding the branch offset to the program counter. Recall the branch offset resides in the *imm* field of the instruction ( $\text{Instr}_{15:0}$ )
  - $\text{PC}' = \text{PC} + 4 + \text{SignImm} \times 4$
  - Multiply by 4 is implemented as shift left by 2 (*PCBranch*)
- If operands are equal and instruction is branch ( $\text{Branch} = 1$ ), then the next instruction comes from *PCBranch*; otherwise, *PCPlus4*



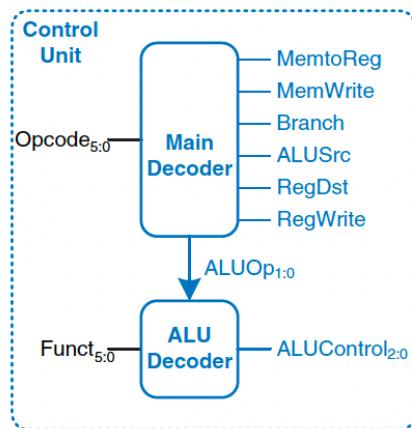
# Single-Cycle Control

- The control unit controls the signals that determine the ALU operation
- Most control information comes from *opcode* but R-Type instructions use the *funct* field too



# Single-Cycle Control

- We simplify the control unit by factoring out two blocks of combinational logic
  - The main decoder computes the output control signals along with a 2-bit *ALUOp* signal
  - The ALU decoder uses *ALUOp* and *funct* to compute the *ALUControl*
  - For our purposes, we will only have to know the meaning of the *ALUOp* signals



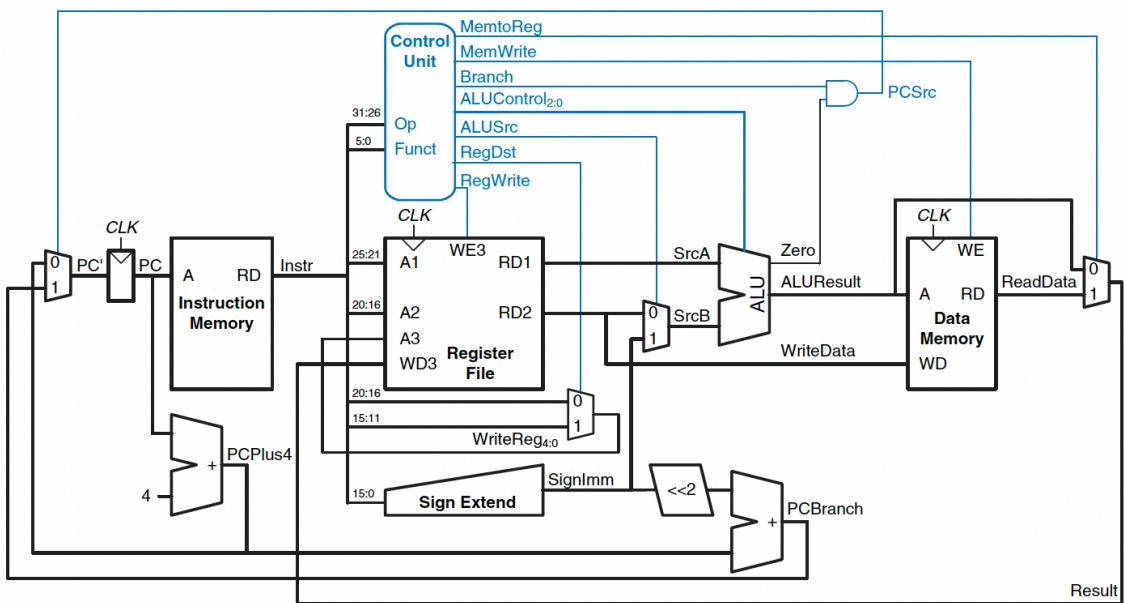
ALUOp	Meaning
00	add
01	subtract
10	look at funct field
11	n/a

ALUOp	Funct	ALUControl
00	X	010 (add)
X1	X	110 (subtract)
1X	100000 (add)	010 (add)
1X	100010 (sub)	110 (subtract)
1X	100100 (and)	000 (and)
1X	100101 (or)	001 (or)
1X	101010 (slt)	111 (set less than)

The *funct* fields for R-type instructions.  
X indicates don't care

# Single-Cycle Control

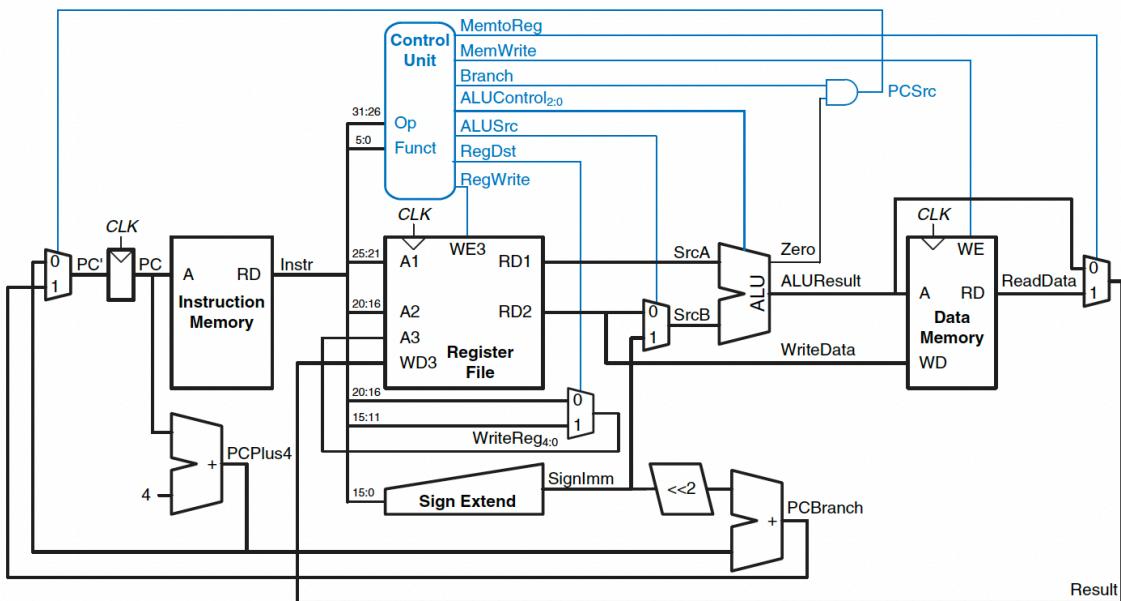
We determine the value for each control signal



Instruction	Opcode	RegWrite	RegDst	ALUSrc	Branch	MemWrite	MemtoReg	ALUOp
R-type	000000							
lw	100011							
sw	101011							
beq	000100							

# Single-Cycle Control

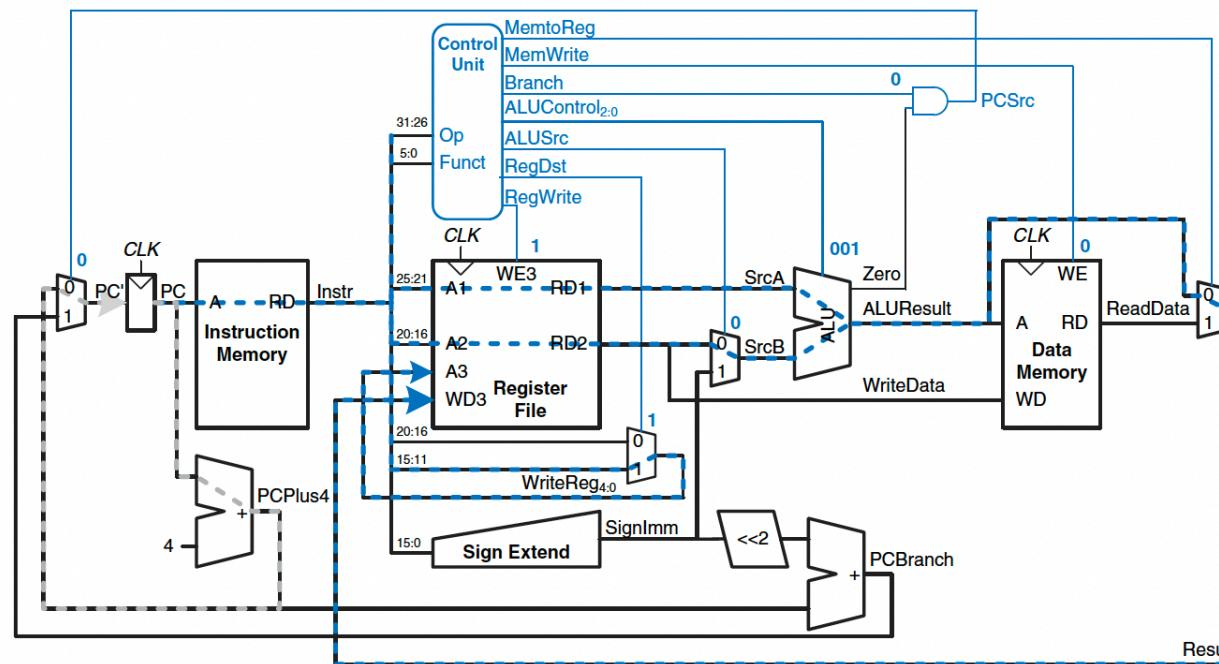
We determine the value for each control signal



Instruction	Opcode	RegWrite	RegDst	ALUSrc	Branch	MemWrite	MemtoReg	ALUOp
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01

# Single-Cycle Example

- As an example, let us track the datapath and control path of the `or` instruction, which R-type
  - The main flow of data through the register file and ALU is represented with a dashed blue line
  - The updating of the PC is shown in dashed grey line

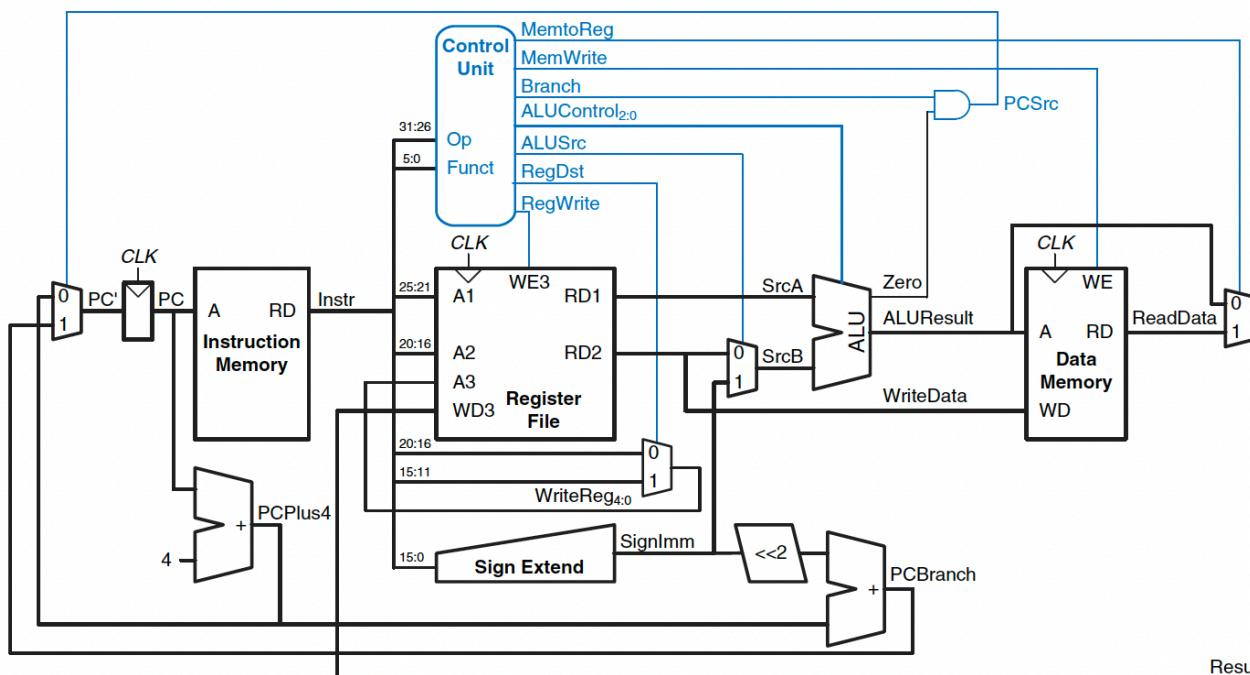


# Extending The Single Cycle CPU

- So far, we have considered a limited set of instructions
- But we can extend the processor for more instructions
- As an example, we will extend the processor with the addi and j instructions

# Extending The Single Cycle CPU

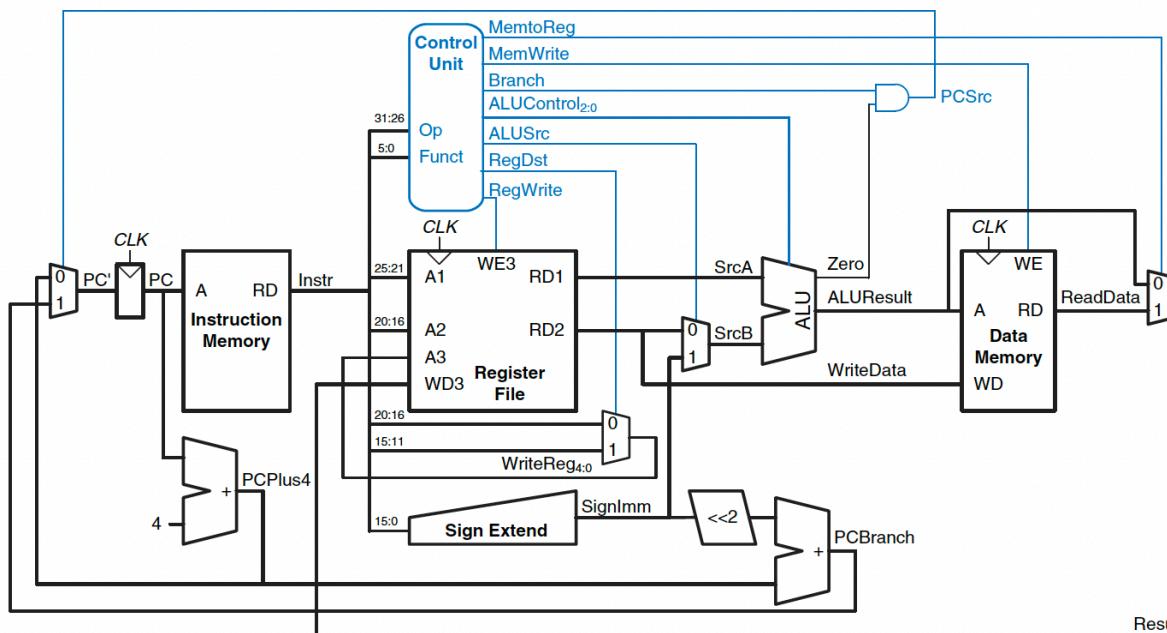
- The `addi` instruction adds the value in a register to an immediate and writes the result to another register.
  - The datapath for `addi` already exists. So, we don't need additional hardware!
  - We only have to determine the correct control signal values



Instruction	Opcode	RegWrite	RegDst	ALUSrc	Branch	MemWrite	MemtoReg	ALUOp
addi	001000							00

# Extending The Single Cycle CPU

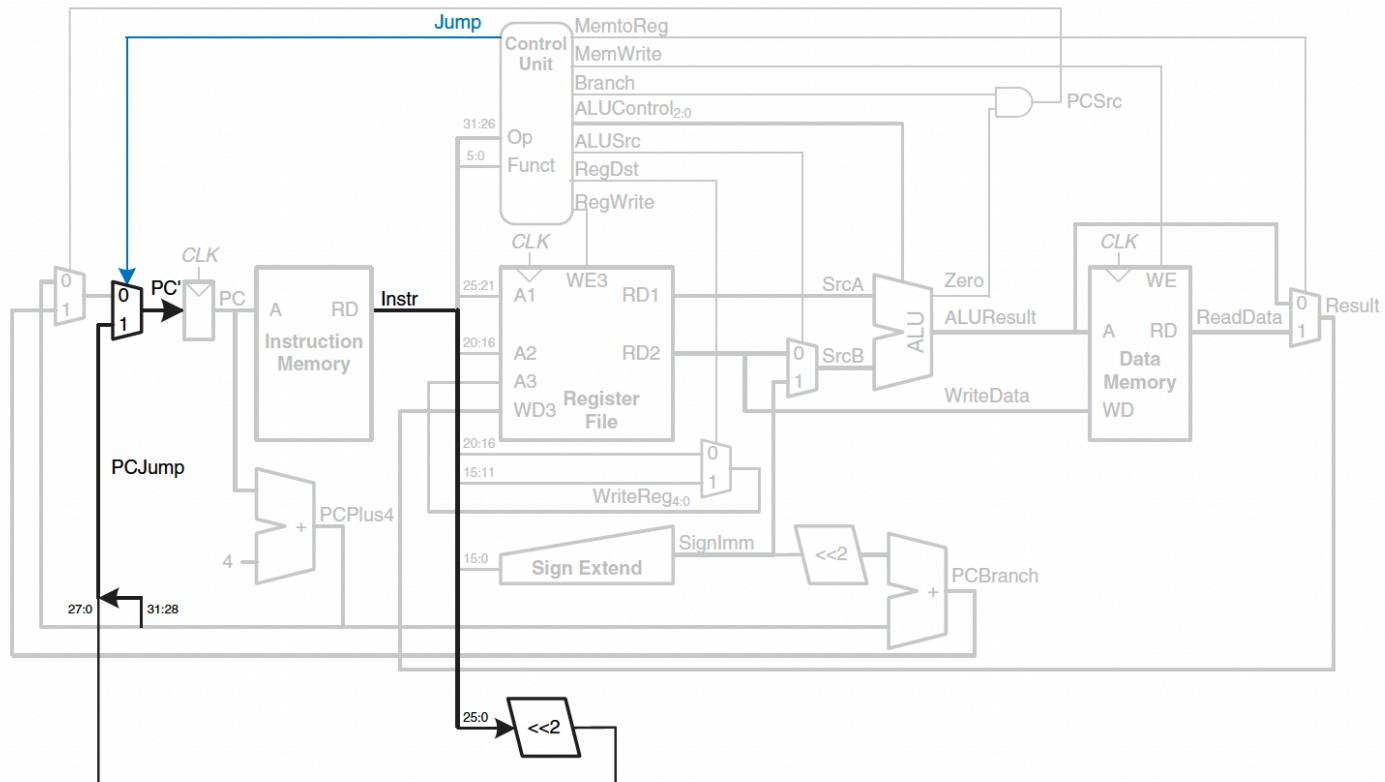
- The addi instruction adds the value in a register to an immediate and writes the result to another register.
  - The datapath for addi already exists. So, we don't need additional hardware!
  - We only have to determine the correct control signal values



Instruction	Opcode	RegWrite	RegDst	ALUSrc	Branch	MemWrite	MemtoReg	ALUOp
addi	001000	1	0	1	0	0	0	00

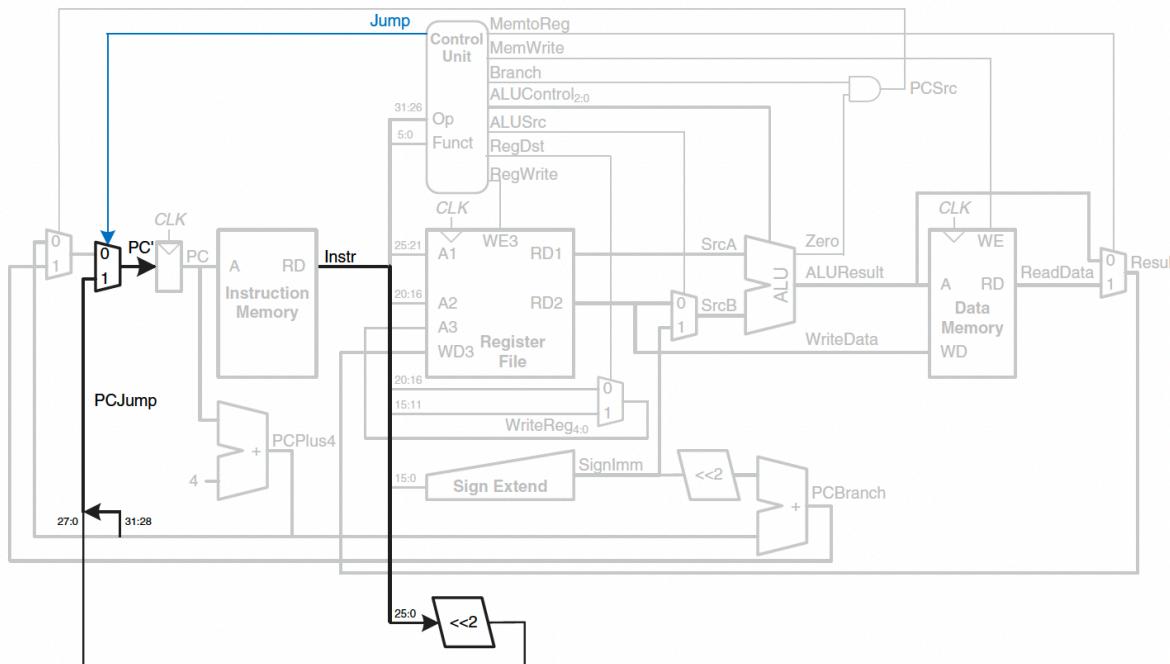
# Extending The Single Cycle CPU

- The jump instruction, writes a new value to the PC
  - The 2 LSBs of the PC are always 0 (because word aligned, i.e., multiple of 4)
  - The next 26 bits are taken from the jump address field of the instruction
  - The upper 4 bits are taken from the old PC
- We don't have hardware to compute PC this way! What do we do?
  - Add hardware, specifically a multiplexer with a control signal *Jump*
  - *Jump = 1* for jump instructions and 0 for others



# Extending The Single Cycle CPU

- Determine the control signals for jump



Instruction	Opcode	RegWrite	RegDst	ALUSrc	Branch	MemWrite	MemtoReg	ALUOp	Jump
j	000010	0	X	X	X	0	X	XX	1

# Points To Note

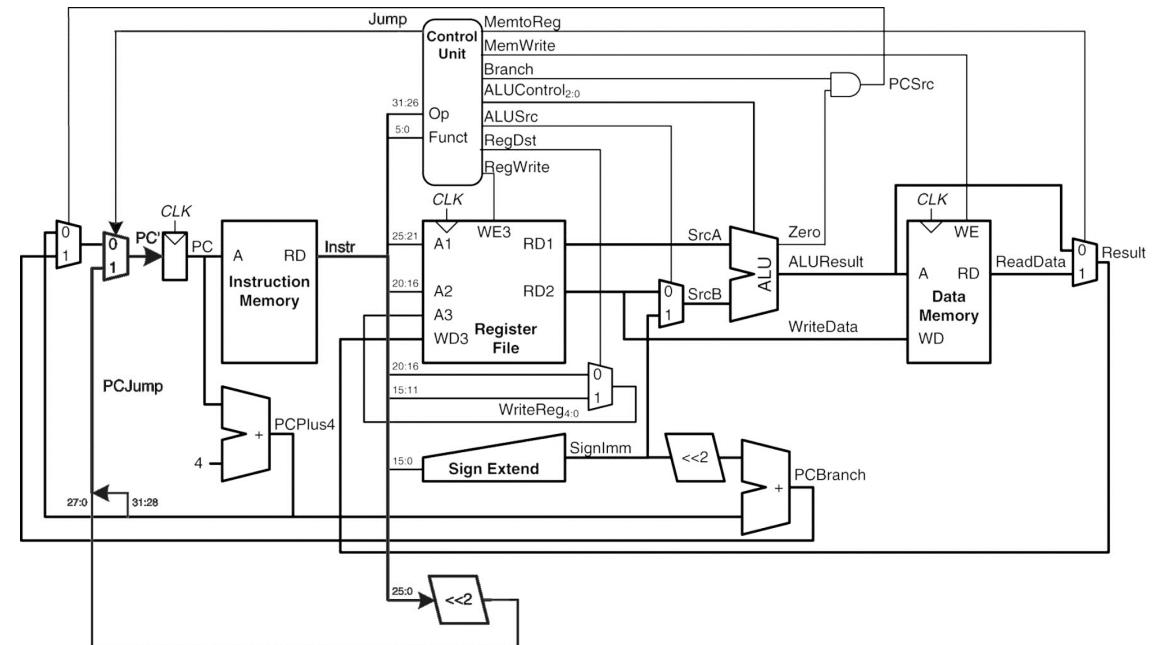
- When we build the hardware for a particular instruction we focus on additional hardware and the control signal values.
- This does not mean the datapath for existing instructions are switched off!
- Signals are always travelling thru the wires.
- It is the engineer/designer's job to control the signals in a way that does not cause the CPU to malfunction.

# Recap

So far, we have built a single cycle processor for a few instructions

Instruction	Opcode	RegWrite	RegDst	ALUSrc	Branch	MemWrite	MemtoReg	ALUOp
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01
addi	001000	1	0	1	0	0	0	00

Instruction	Opcode	RegWrite	RegDst	ALUSrc	Branch	MemWrite	MemtoReg	ALUOp	Jump
j	000010	0	X	X	X	0	X	XX	1

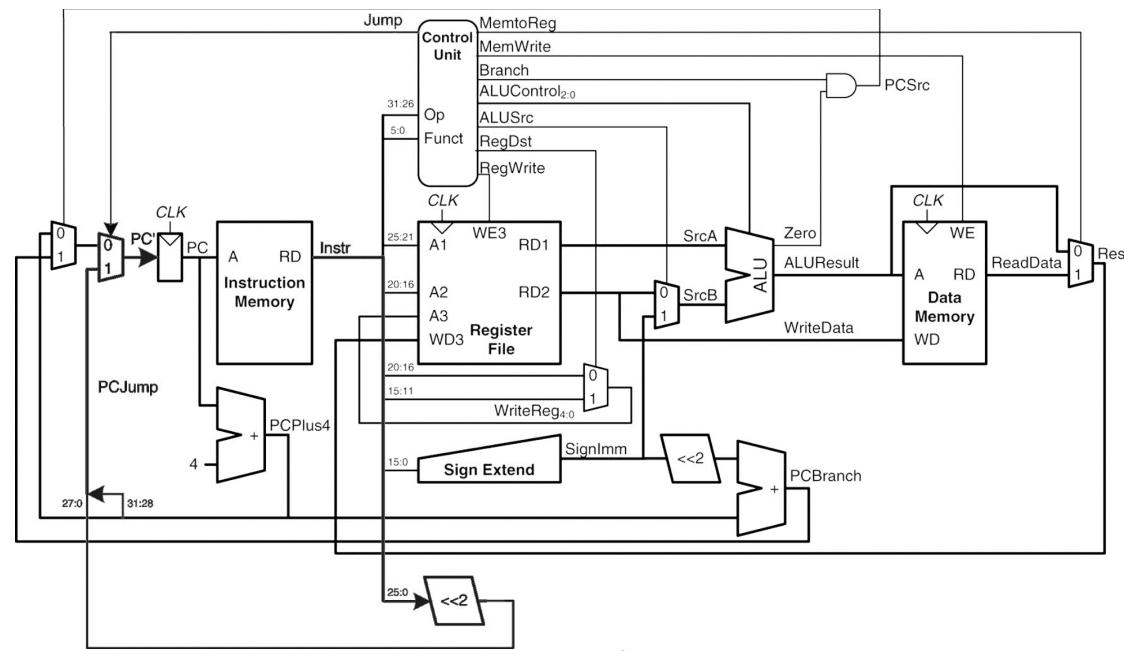


# Performance Analysis

- Recall, we mentioned that different microarchitectures have different performance and cost tradeoffs
- There are many ways to measure performance; marketing departments infamously choose the measure that makes their computer look fastest!
- The only gimmick-free way to measure performance is to measure the execution of a program of interest to us or benchmarks
  - Execution Time = (# instructions) (cycles/instruction) (seconds/cycle)
- The no. of cycles per instruction is called the CPI
  - For single-cycle CPUs, CPI = 1
- The no. of seconds per cycle is called the clock period,  $T_c$ 
  - Determined by critical path thru the logic on the processor

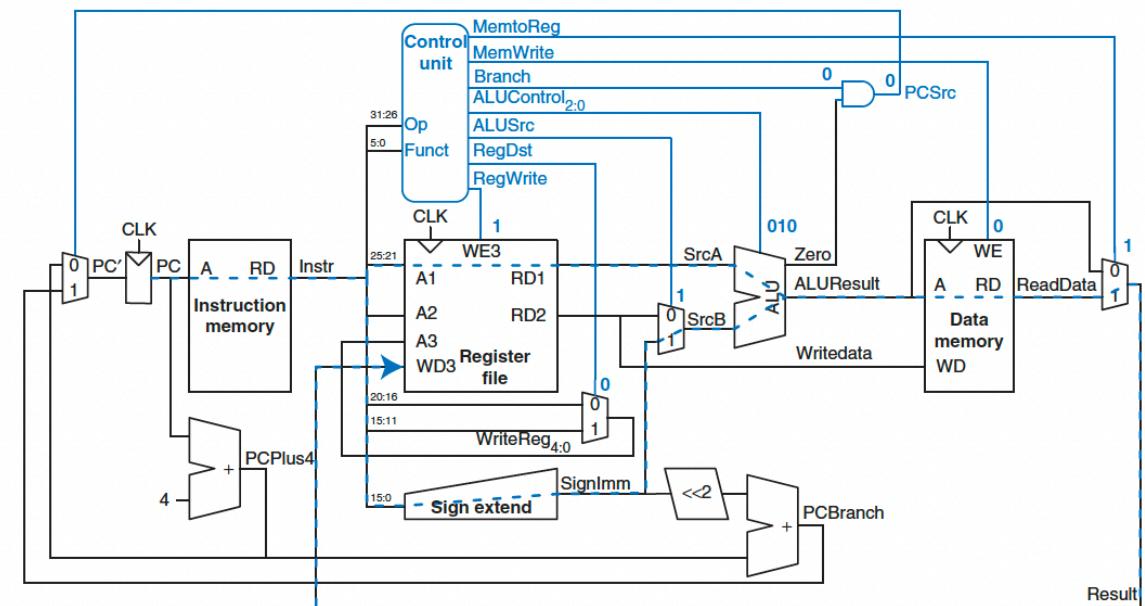
# Performance Analysis

- Suppose we want to determine the clock period of the single cycle processor that we have been implementing
- We need to compute the delay along the critical path of the slowest instruction, which happens to be the `lw` instruction



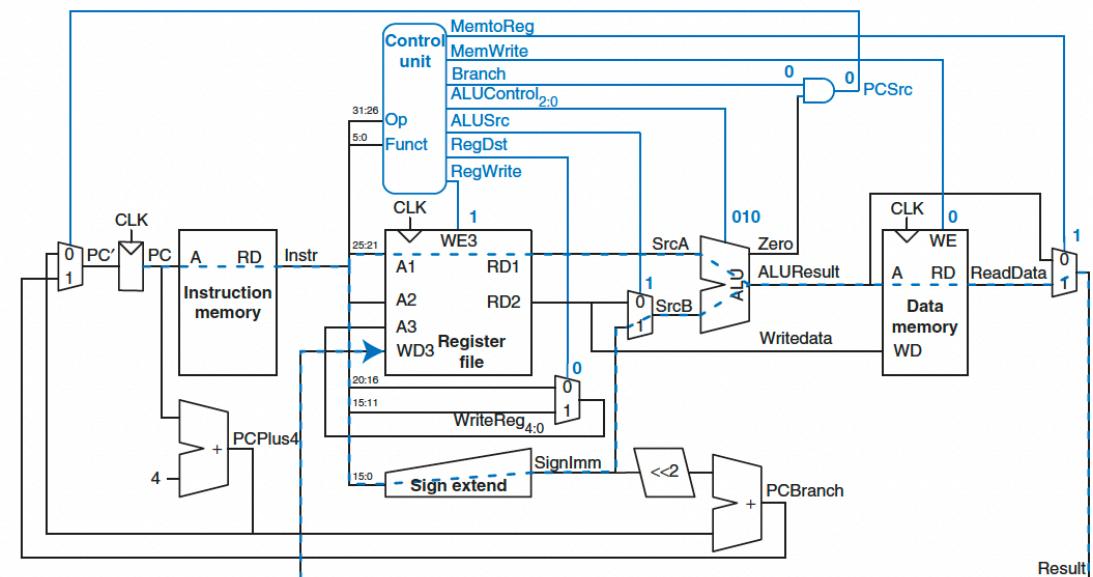
# Critical Path of $lw$ Instruction

- It starts with the PC loading a new address at the rising edge of the clock
- The instruction memory reads the next instruction
- The register file reads  $SrcA$
- While register file is read, the immediate is sign extended and is available at the  $ALUSrc$  MUX to determine  $SrcB$
- The ALU adds  $SrcA$  and  $SrcB$  to find effective address
- The data memory reads from this address
- The *MemtoReg* MUX selects *ReadData*
- Finally, *Result* setup at the register file before the rising edge of the clock



# Critical Path of $lw$ Instruction

- It starts with the PC loading a new address at the rising edge of the clock ( $t_{pcq\_PC}$ )
- The instruction memory reads the next instruction ( $t_{mem}$ )
- The register file reads  $SrcA$  ( $t_{RFread}$ )
- While register file is read, the immediate is sign extended and is available at the  $ALUSrc$  MUX to determine  $SrcB$  ( $t_{sext} + t_{mux}$ )
- The ALU adds  $SrcA$  and  $SrcB$  to find effective address ( $t_{ALU}$ )
- The data memory reads from this address ( $t_{mem}$ )
- The  $MemtoReg$  MUX selects  $ReadData$  ( $t_{mux}$ )
- Finally,  $Result$  setup at the register file before the rising edge of the clock ( $t_{RFsetup}$ )



# Critical Path of 1w Instruction

- In most implementations, the ALU, memory, register file accesses are slower than other operations.
- Therefore, the cycle time simplifies:
  - $T_c = t_{pcq\_PC} + t_{mem} + t_{RFread} + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$
- We don't consider  $t_{sext} + t_{mux}$  and the time to compute the new PC. Why?
  - Because they happen in parallel with the other slower operations
  - Slower operations dominate Now, recall there are other instructions for which we can come up with a formula
- But they have shorter cycle times. E.g., R-type instructions do not access memory
- The clock period must be long enough to accommodate the slowest instruction

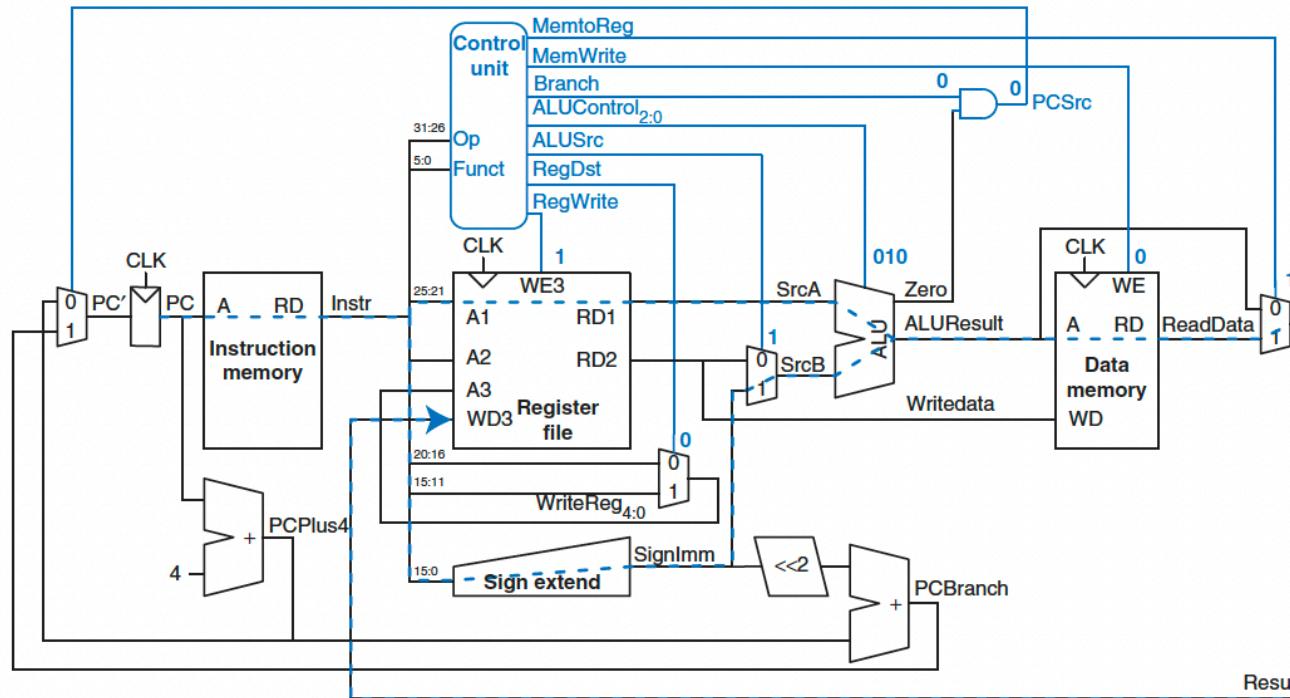
# Critical Path of 1w Instruction

- Generally, you can assume that the delays for each of the circuit elements in our processor is known
- Hence, we can concretize the clock cycle now that we know the values

Element	Parameter	Delay (ps)
register clk-to-Q	$t_{pcq}$	30
register setup	$t_{\text{setup}}$	20
multiplexer	$t_{\text{mux}}$	25
ALU	$t_{\text{ALU}}$	200
memory read	$t_{\text{mem}}$	250
register file read	$t_{RF\text{read}}$	150
register file setup	$t_{RF\text{setup}}$	20

# Critical Path of lw Instruction

- We can use the formula we came up with.



$$T_c = 30 + 250 + 150 + 200 + 250 + 25 + 20 \\ = 925 \text{ ps}$$

$$T_c = 30 + 250 + 150 + 200 + 250 + 25 + 20 \\ = 925 \text{ ps}$$

# Critical Path of $l_w$ Instruction

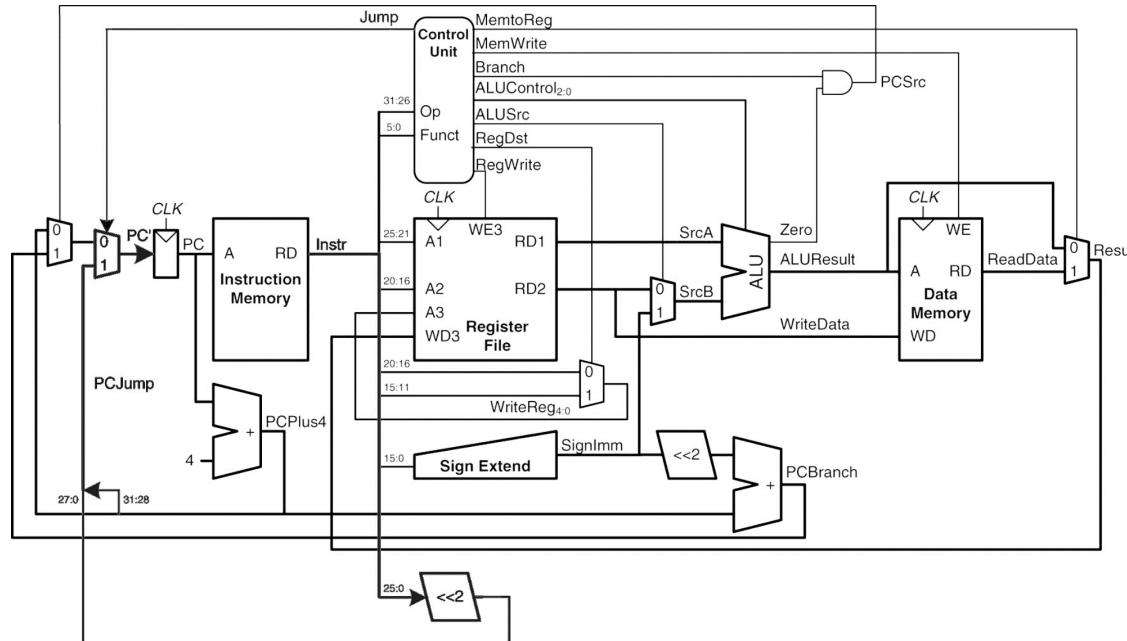
- Suppose we want to compute the execution time ( $T_1$ ) of a program with 100 billion instructions
- Assuming that  $l_w$  is the slowest instruction, we get:

$$T_c = 30 + 250 + 150 + 200 + 250 + 25 + 20 = 925 \text{ ps}$$

$$T_1 = (100 * 10^9)(1 \text{ cycle/instruction})(925 * 10^{-12} \text{s/cycle}) = 92.5 \text{ seconds}$$

# Critical Path of R-Type Instruction

- Let's compute the critical path delay (propagation delay) for R-type instructions in the datapath



$t_{pcq}$	I-Mem Read	Adder	MUX	ALU	RegFile Read	RegFile Write	D-Mem Read	D-Mem Write	SignExt	Shifter
30	40	5	3	20	30	45	50	40	4	6

# Critical Path of R-Type Instruction

- Determine he cycle time formula:
  - $T_c = t_{pcq\_PC} + t_{mem} + t_{RFread} + t_{mux} + t_{ALU} + t_{mux} + t_{RFwrite}$
- Ignore  $t_{sext} + t_{mux}$  and the time to compute the new PC because negligible
- $T_c = 30 + 40 + 30 + 3 + 20 + 3 + 45 = 171$

# Custom Instructions

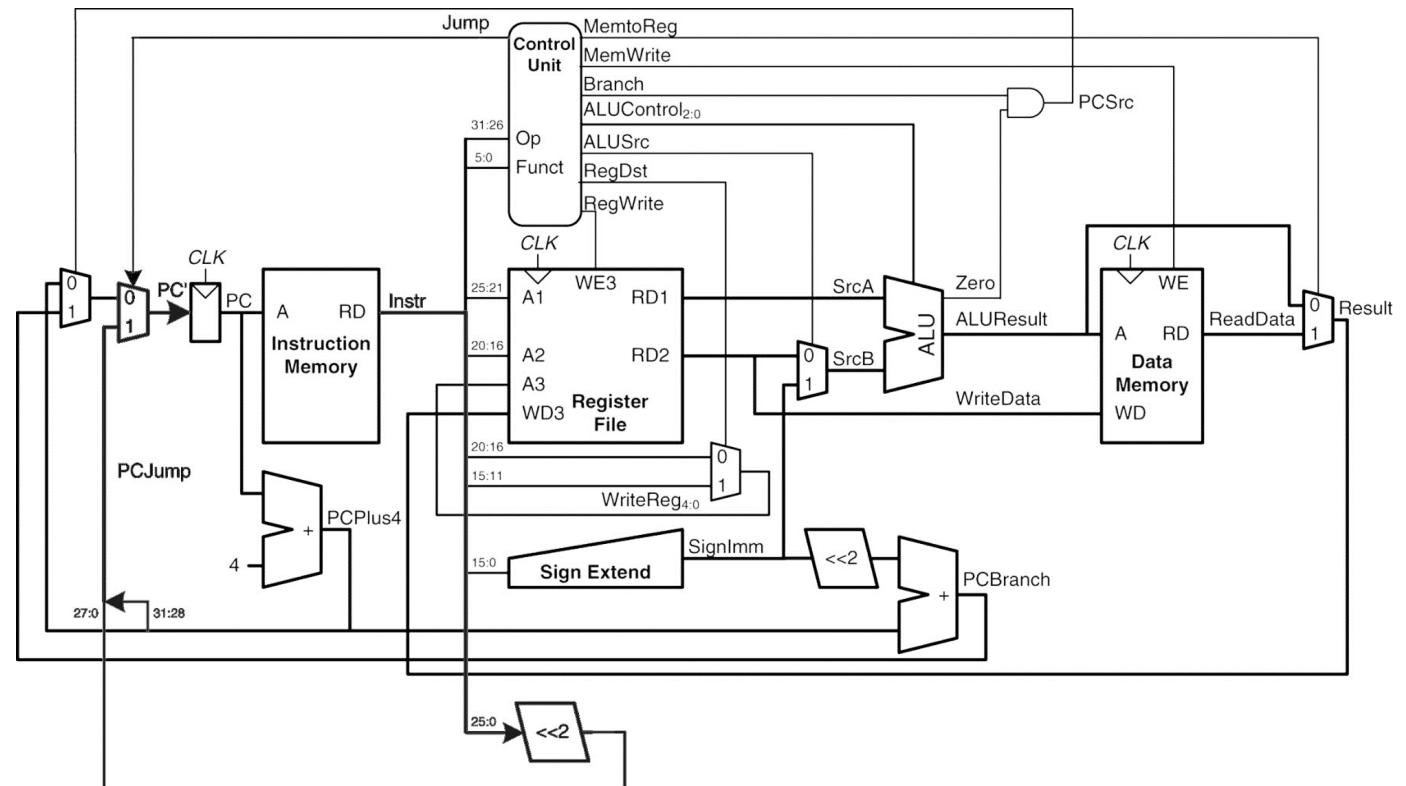
- We will now extend our single cycle datapath with an instruction that we have invented!
- `xcond $rs ,$rt, immediate` is an I-type instruction defined as follows:

```
if Reg[$rt] >= 0 then
    Mem[SignImm] = Reg[$rs]
else
    do nothing
PC' = PC + 4
```
- How should we implement this?
  - Notice that `xcond` is a conditional memory write
  - This means that write to memory happens only when the condition true
  - Hence, we need to optionally switch on or switch off the `MemWrite` signal
  - We also need to specify the write address and the data that will be written
- Hardware assumptions:
  - Assume that we can add at most one logical gate and three 2:1 multiplexers

# Custom Instructions

- `xcond $rs , $rt, immediate` is an I-type instruction defined as follows:

```
if Reg[$rt] >= 0 then
    Mem[SignImm] = Reg[$rs]
else
    do nothing
PC' = PC + 4
```

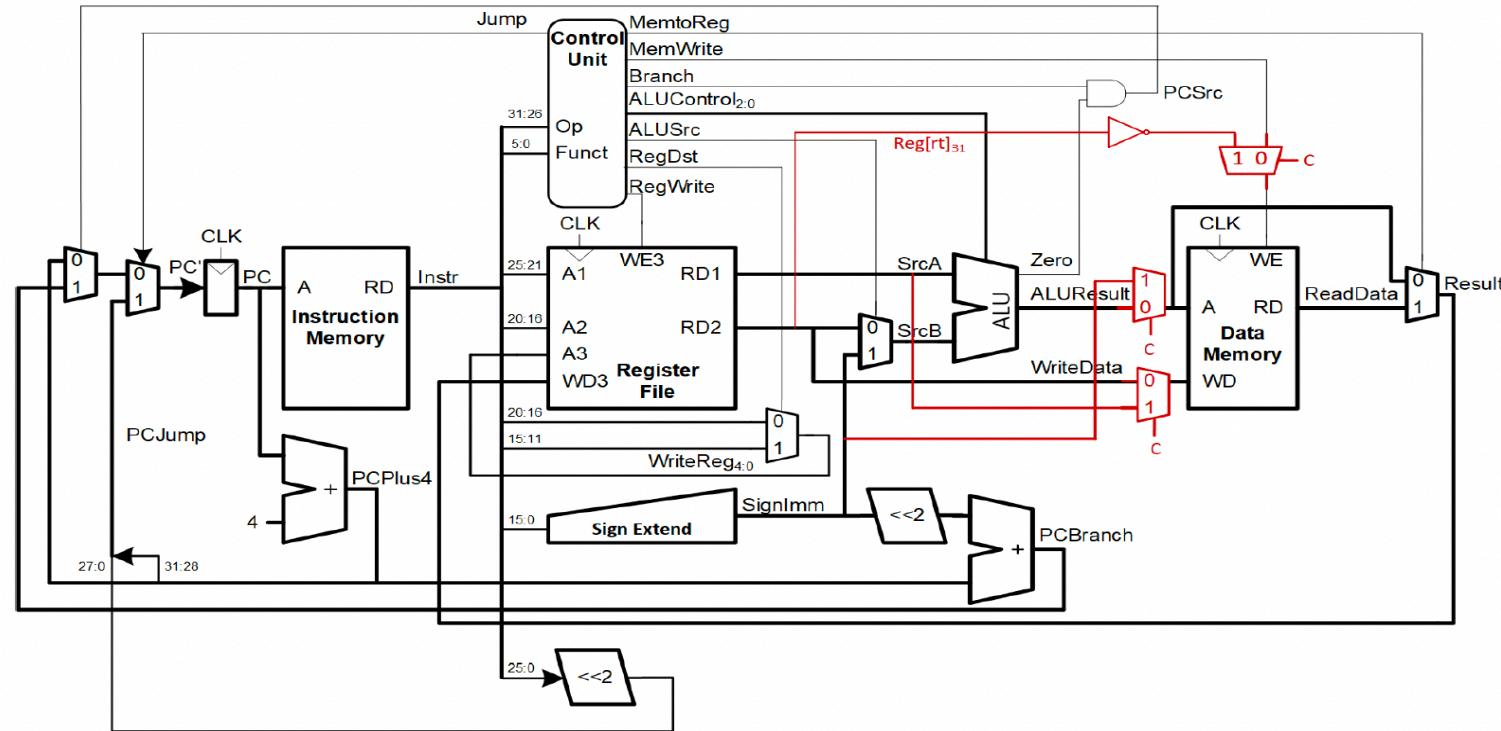


# Custom Instructions

- xcond \$rs ,\$rt, immediate is an I-type instruction defined as follows:

```

if Reg[rt] >= 0 then
    Mem[SignImm] = Reg[$rs]
else
    do nothing
PC' = PC + 4
  
```

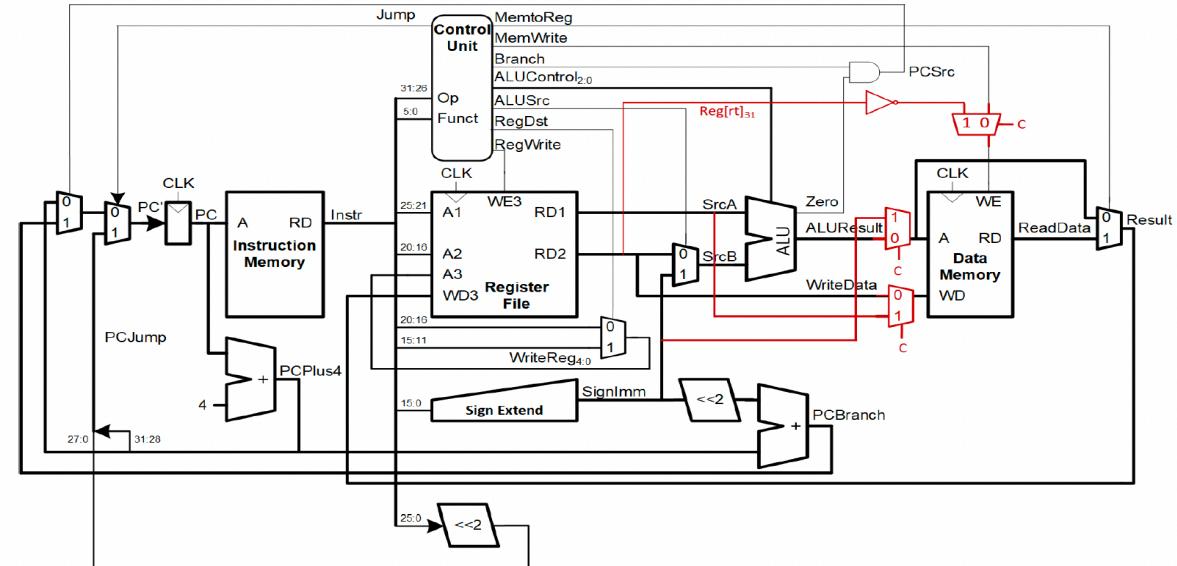


# Custom Instructions

- `xcond $rs ,$rt, immediate` is an I-type instruction defined as follows:

```

if Reg[rt] >= 0 then
    Mem[SignImm] = Reg[$rs]
else
    do nothing
PC' = PC + 4
  
```



Instruction	RegDst	ALUSrc	MemtoReg	RegWrite	MemWrite	Branch	ALUOp	Jump	C
R-type	1	0	0	1	0	0	10	0	0
lw	0	1	1	1	0	0	00	0	0
sw	X	1	X	0	1	0	00	0	0
addi	0	1	0	1	0	0	00	0	0
beq	X	0	X	0	0	1	01	0	0
j	X	X	X	0	0	X	XX	1	0
xcond									

# Custom Instructions

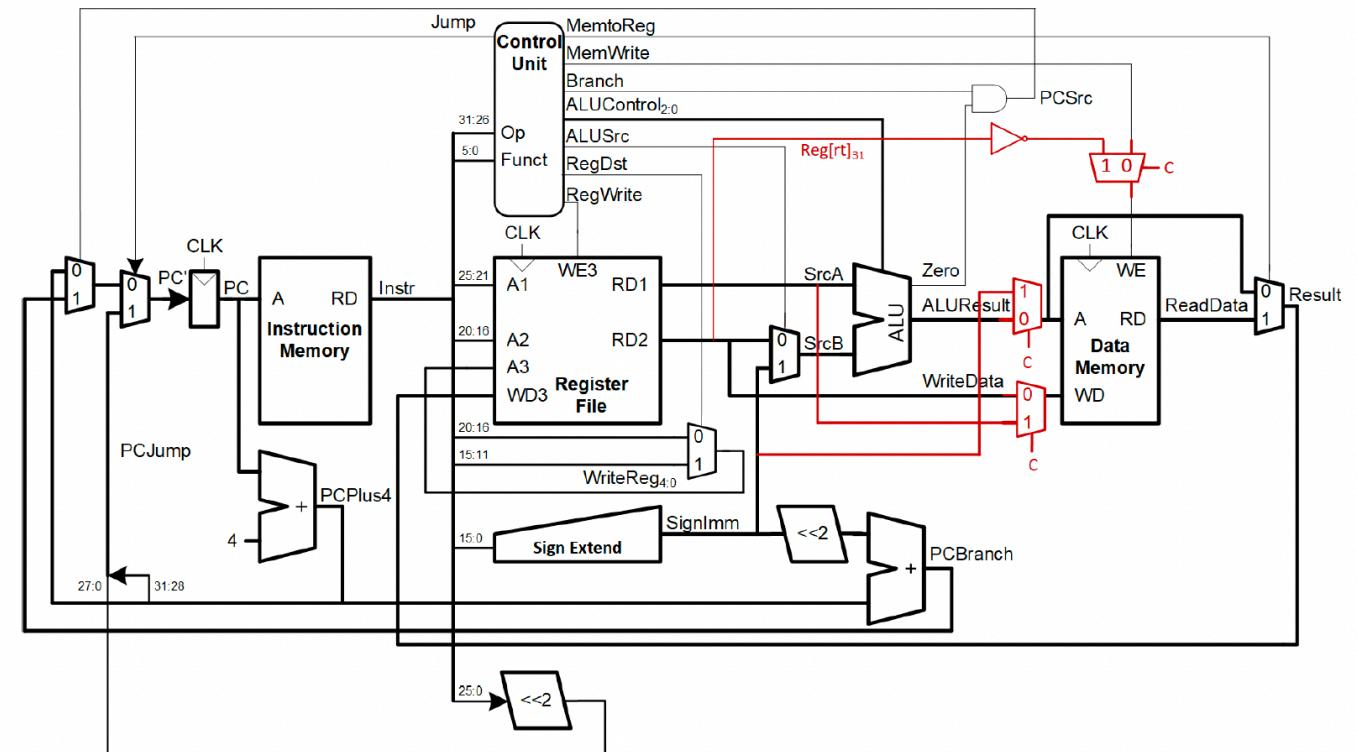
- xcond \$rs , \$rt, immediate is an I-type instruction defined as follows:

```

if Reg[rt] >= 0 then
    Mem[SignImm] = Reg[$rs]
else
    do nothing
PC' = PC + 4

```

Type	RegDst	ALUSrc	MemtoReg	RegWrite	MemWrite	Branch	ALUOp <sub>1:0</sub>	Jump	c
R-type	1	0	0	1	0	0	10	0	0
lw	0	1	1	1	0	0	00	0	0
sw	X	1	X	0	1	0	00	0	0
addi	0	1	0	1	0	0	00	0	0
beq	X	0	X	0	0	1	01	0	0
j	X	X	X	0	0	X	XX	1	0
xcond	X	X	X	0	X	0	XX	0	1



# Rules For Adding Instructions

- Determine the type of instruction (R-type, I-type, or J-type).
- Compare the new instructions with the existing instructions in the datapath.
- Ask yourself if the existing datapath is enough or additional hardware is required.
- If existing datapath is enough, then analyze the control signal values needed to implement the instruction.
- If additional hardware is needed, then try and use minimal hardware.
- Ensure that the new datapath still works for existing instructions!

# Custom Instructions

- Imagine a MIPS instruction that would take a branch if the value in `rs` were not a multiple of 4: `bnmul4 rs, label`

```
if Reg[rs] is NOT a multiple of 4 then
    PC' = PC+4 + SignImm×4
else
    PC' = PC+4
```
- How could we implement this? What hardware and logic would we need to add to the single-cycle datapath and control unit?
- The binary representation of an integer (positive or negative) that is a multiple of four will end in 00
- So, we will branch if either or both of the two least significant bits are 1s:  
`Reg[rs][0] = 1 or Reg[rs][1] = 1`

# Custom Instructions

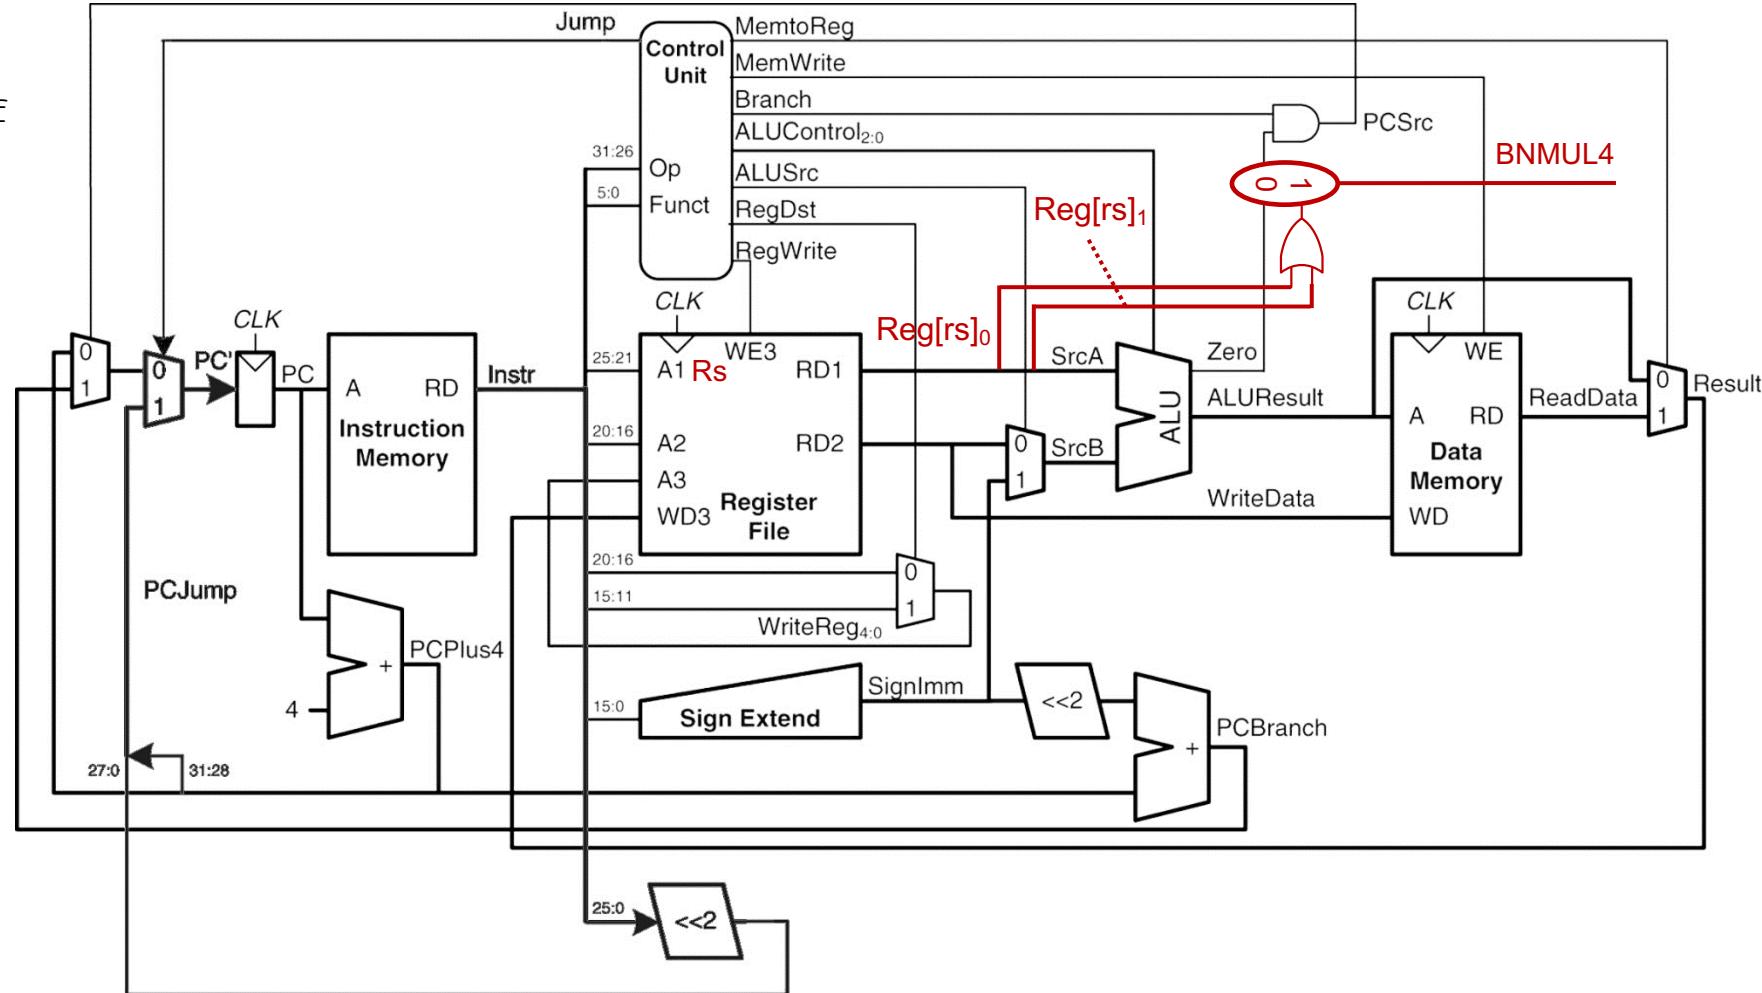
bnmul4 rs, label:

if Reg[rs] is NOT a multiple of  
then

$PC' = PC + 4 + \text{SignImm} \times 4$

else

$PC' = PC + 4$



# Custom Instructions

- Control signals

